

Learning Tractable Graphical Models: Latent Trees

Furong Huang


U.C. Irvine

Joint work with Anima Anandkumar and U.N. Niranjan.

High-Dimensional Graphical Modeling


Modeling Conditional Independencies through Graphs

- $X_u \perp\!\!\!\perp X_v | X_S$.
- Learning and inference are NP-hard.



Tractable Models: Tree Models


- Efficient inference using belief propagation



Walk-up: Learning Tree Models


Data processing inequality for Markov chains

$$I(X_1; X_3) \leq I(X_1; X_2), I(X_2; X_3).$$



Tree Structure Estimation (Chow and Liu '68)


- **MLE**: Max-weight tree with estimated mutual information weights



Walk-up: Learning Tree Models


Data processing inequality for Markov chains

$$I(X_1; X_3) \leq I(X_1; X_2), I(X_2; X_3).$$



Tree Structure Estimation (Chow and Liu '68)


- **MLE**: Max-weight tree with estimated mutual information weights
- **Pairwise** statistics suffice



Walk-up: Learning Tree Models

Data processing inequality for Markov chains


$$I(X_1; X_3) \leq I(X_1; X_2), I(X_2; X_3).$$



Tree Structure Estimation (Chow and Liu '68)

- **MLE**: Max-weight tree with estimated mutual information weights
- **Pairwise** statistics suffice
- n samples and p nodes


Sample complexity: $\frac{\log p}{n} = O(1).$



Learning Tractable Graphical Models

Tractable Models: Tree Models


- Efficient inference using **belief propagation**
- MLE is easy to compute.
- Tree models are highly restrictive.



Learning Tractable Graphical Models


Tractable Models: Tree Models

- Efficient inference using **belief propagation**
- MLE is easy to compute.
- Tree models are highly restrictive.




Latent tree graphical models

- Tree models with hidden variables.
- **Number** and **location** of hidden variables unknown.



Application: Hierarchical Topic Modeling

- Data: Word co-occurrences.
- Graph: Topic-word structure.



Application of Latent Trees: Object Recognition


- **Challenge:** Succinct representation of large-scale data
 - ▶ **Input:** ~ 100 object categories, ~ 4000 training images
 - ▶ **Goal:** learn $\sim 2^{100}$ co-occurrence probabilities
 - **Solution:** Latent tree graphical models
-



“Context Models and Out-of-context Objects,” M. J. Choi, A. Torralba, and A. S. Willsky, Pattern Recognition Letters, 2012.

Application of Latent Trees: Object Recognition

- **Challenge:** Succinct representation of large-scale data
 - ▶ **Input:** ~ 100 object categories, ~ 4000 training images
 - ▶ **Goal:** learn $\sim 2^{100}$ co-occurrence probabilities
- **Solution:** Latent tree graphical models




In this talk: learning latent tree models and tree mixtures.

"Context Models and Out-of-context Objects," M. J. Choi, A. Torralba, and A. S. Willsky, Pattern Recognition Letters, 2012.

Summary of Results

Latent Tree Models


- **Number** of hidden variables and **location** unknown
- Integrated structure and parameter estimation.
- **Local** learning with **global** consistency guarantees.



Outline

- 1 Introduction
- 2 Tests for Structure Learning**
- 3 Parameter Learning through Tensor Methods
- 4 Integrating Structure and Parameter Learning
- 5 Conclusion


Learning Latent Tree Graphical Models



Linear Multivariate Models


- Conditional independence w.r.t tree
- **Categorical** k -state hidden variables.
- **Multivariate** d -dimensional observed variables. $k \leq d$.
- When y is nbr. of h , $\mathbb{E}[y|h] = Ah$.
- Includes discrete, Poisson and Gaussian models, Gaussian mixtures etc.

Additive Tree Distances



Information Distances $[d_{i,j}]$ for Tree Models


Additive Tree Distances



Information Distances $[d_{i,j}]$ for Tree Models

Gaussian scalar: $d_{ij} := -\log |\rho_{ij}|$.

Additive Tree Distances




Information Distances $[d_{i,j}]$ for Tree Models

Gaussian scalar: $d_{ij} := -\log |\rho_{ij}|$.

Linear multivariate models: $d_{ij} := -\log \frac{\prod_{\sigma_i \neq 0} \sigma(\mathbb{E}[y_i y_j^\top])}{\sqrt{\det \mathbb{E}[y_i y_i^\top] \det \mathbb{E}[y_j y_j^\top]}}$.

Additive Tree Distances



Information Distances $[d_{i,j}]$ for Tree Models


Gaussian scalar: $d_{ij} := -\log |\rho_{ij}|$.

Linear multivariate models: $d_{ij} := -\log \frac{\prod_{\sigma_i \neq 0} \sigma(\mathbb{E}[y_i y_j^\top])}{\sqrt{\det \mathbb{E}[y_i y_i^\top] \det \mathbb{E}[y_j y_j^\top]}}$.

$[d_{i,j}]$ is an additive tree metric:

$$d_{k,l} = \sum_{(i,j) \in \text{Path}(k,l;E)} d_{i,j}.$$

Additive Tree Distances



Information Distances $[d_{i,j}]$ for Tree Models


Gaussian scalar: $d_{ij} := -\log |\rho_{ij}|$.

Linear multivariate models: $d_{ij} := -\log \frac{\prod_{\sigma_i \neq 0} \sigma(\mathbb{E}[y_i y_j^\top])}{\sqrt{\det \mathbb{E}[y_i y_i^\top] \det \mathbb{E}[y_j y_j^\top]}}$.

$[d_{i,j}]$ is an additive tree metric:

$$d_{k,l} = \sum_{(i,j) \in \text{Path}(k,l;E)} d_{i,j}.$$

Additive Tree Distances



Information Distances $[d_{i,j}]$ for Tree Models


Gaussian scalar: $d_{ij} := -\log |\rho_{ij}|$.

Linear multivariate models: $d_{ij} := -\log \frac{\prod_{\sigma_i \neq 0} \sigma(\mathbb{E}[y_i y_j^\top])}{\sqrt{\det \mathbb{E}[y_i y_i^\top] \det \mathbb{E}[y_j y_j^\top]}}$.

$[d_{i,j}]$ is an additive tree metric:

$$d_{k,l} = \sum_{(i,j) \in \text{Path}(k,l;E)} d_{i,j}.$$

Additive Tree Distances



Information Distances $[d_{i,j}]$ for Tree Models

Gaussian scalar: $d_{ij} := -\log |\rho_{ij}|$.

Linear multivariate models: $d_{ij} := -\log \frac{\prod_{\sigma_i \neq 0} \sigma(\mathbb{E}[y_i y_j^\top])}{\sqrt{\det \mathbb{E}[y_i y_i^\top] \det \mathbb{E}[y_j y_j^\top]}}$.

$[d_{i,j}]$ is an additive tree metric:

$$d_{k,l} = \sum_{(i,j) \in \text{Path}(k,l;E)} d_{i,j}.$$


Learning latent tree using $[\hat{d}_{i,j}]$

Siblings Test Based on Information Distances

Exact Statistics: Distances $[d_{i,j}]$

Let $\Phi_{ijk} := d_{i,k} - d_{j,k}$.

- $-d_{i,j} < \Phi_{ijk} = \Phi_{ijk'} < d_{i,j} \quad \forall k, k' \neq i, j, \iff i, j$ leaves with common parent
- $\Phi_{ijk} = d_{i,j}, \quad \forall k \neq i, j, \iff i$ is a leaf and j is its parent.



Sample Statistics: ML Estimates $[\hat{d}_{i,j}]$


Use only short distances: $d_{i,k}, d_{j,k} < \tau$, Relax equality relationships

Siblings Test Based on Information Distances

Exact Statistics: Distances $[d_{i,j}]$

Let $\Phi_{ijk} := d_{i,k} - d_{j,k}$.

- $-d_{i,j} < \Phi_{ijk} = \Phi_{ijk'} < d_{i,j} \quad \forall k, k' \neq i, j, \iff i, j$ leaves with common parent
- $\Phi_{ijk} = d_{i,j}, \quad \forall k \neq i, j, \iff i$ is a leaf and j is its parent.



Sample Statistics: ML Estimates $[\hat{d}_{i,j}]$


Use only short distances: $d_{i,k}, d_{j,k} < \tau$, Relax equality relationships

Siblings Test Based on Information Distances

Exact Statistics: Distances $[d_{i,j}]$

Let $\Phi_{ijk} := d_{i,k} - d_{j,k}$.

- $-d_{i,j} < \Phi_{ijk} = \Phi_{ijk'} < d_{i,j} \quad \forall k, k' \neq i, j, \iff i, j$ leaves with common parent
- $\Phi_{ijk} = d_{i,j}, \quad \forall k \neq i, j, \iff i$ is a leaf and j is its parent.



Sample Statistics: ML Estimates $[\hat{d}_{i,j}]$


Use only short distances: $d_{i,k}, d_{j,k} < \tau$, Relax equality relationships

Siblings Test Based on Information Distances

Exact Statistics: Distances $[d_{i,j}]$

Let $\Phi_{ijk} := d_{i,k} - d_{j,k}$.

- $-d_{i,j} < \Phi_{ijk} = \Phi_{ijk'} < d_{i,j} \quad \forall k, k' \neq i, j, \iff i, j$ leaves with common parent
- $\Phi_{ijk} = d_{i,j}, \quad \forall k \neq i, j, \iff i$ is a leaf and j is its parent.



Sample Statistics: ML Estimates $[\hat{d}_{i,j}]$


Use only short distances: $d_{i,k}, d_{j,k} < \tau$, Relax equality relationships

Siblings Test Based on Information Distances

Exact Statistics: Distances $[d_{i,j}]$

Let $\Phi_{ijk} := d_{i,k} - d_{j,k}$.

- $-d_{i,j} < \Phi_{ijk} = \Phi_{ijk'} < d_{i,j} \quad \forall k, k' \neq i, j, \iff i, j$ leaves with common parent
- $\Phi_{ijk} = d_{i,j}, \quad \forall k \neq i, j, \iff i$ is a leaf and j is its parent.



Sample Statistics: ML Estimates $[\hat{d}_{i,j}]$


Use only short distances: $d_{i,k}, d_{j,k} < \tau$, Relax equality relationships

Siblings Test Based on Information Distances

Exact Statistics: Distances $[d_{i,j}]$

Let $\Phi_{ijk} := d_{i,k} - d_{j,k}$.

- $-d_{i,j} < \Phi_{ijk} = \Phi_{ijk'} < d_{i,j} \quad \forall k, k' \neq i, j, \iff i, j$ leaves with common parent
- $\Phi_{ijk} = d_{i,j}, \quad \forall k \neq i, j, \iff i$ is a leaf and j is its parent.




Sample Statistics: ML Estimates $[\hat{d}_{i,j}]$

Use only short distances: $d_{i,k}, d_{j,k} < \tau$, Relax equality relationships

Recursive Grouping

Recursive Grouping Algorithm (Choi, Tan, Anandkumar, Willsky)

- Sibling test and remove leaves
- Build tree from bottom up





- **Consistent** structure estimation.
- **Serial** method, high computational complexity.

Recursive Grouping

Recursive Grouping Algorithm (Choi, Tan, Anandkumar, Willsky)

- Sibling test and remove leaves
- Build tree from bottom up





- **Consistent** structure estimation.
- **Serial** method, high computational complexity.

Recursive Grouping

Recursive Grouping Algorithm (Choi, Tan, Anandkumar, Willsky)

- Sibling test and remove leaves
- Build tree from bottom up





- **Consistent** structure estimation.
- **Serial** method, high computational complexity.

Recursive Grouping

Recursive Grouping Algorithm (Choi, Tan, Anandkumar, Willsky)

- Sibling test and remove leaves
- Build tree from bottom up





- **Consistent** structure estimation.
- **Serial** method, high computational complexity.

Recursive Grouping

Recursive Grouping Algorithm (Choi, Tan, Anandkumar, Willsky)

- Sibling test and remove leaves
- Build tree from bottom up





- **Consistent** structure estimation.
- **Serial** method, high computational complexity.

Recursive Grouping

Recursive Grouping Algorithm (Choi, Tan, Anandkumar, Willsky)

- Sibling test and remove leaves
- Build tree from bottom up



- **Consistent** structure estimation.
- **Serial** method, high computational complexity.


Outline

- 1 Introduction
- 2 Tests for Structure Learning
- 3 Parameter Learning through Tensor Methods**
- 4 Integrating Structure and Parameter Learning
- 5 Conclusion

Overview of Proposed Parameter Learning Method

Toy Model: 3-star

- Linear multivariate model.
- $A_{x_i|h}^r := \mathbb{E}(x_i|h = e_r)$. and $\lambda_r := \mathbb{P}[h = e_r]$.



$$\mathbb{E}(x_1 \otimes x_2 \otimes x_3) = \sum_{r=1}^k \lambda_r A_{x_1|h}^r \otimes A_{x_2|h}^r \otimes A_{x_3|h}^r.$$


Guaranteed Recovery through Tensor Decomposition

- Transition matrices $A_{x_i|h}$ have full column rank.
- Linear algebraic operations: **SVD** and **tensor power iterations**.

“Tensor Decompositions for Learning Latent Variable Models” by A. Anandkumar, R. Ge, D. Hsu, S.M. Kakade and M. Telgarsky. Preprint, October 2012.

Overview of Tensor Decomposition Technique

- Let $a_r = \mathbb{E}(x_i | h = e_r)$ for all i and $\lambda_r := \mathbb{P}[h = e_r]$.
- $M_3 = \mathbb{E}[x_1 \otimes x_2 \otimes x_3] = \sum_{i=1}^k \lambda_i a_i^{\otimes 3}$.




Intuition: if a_i are orthogonal

- $M_3(I, a_1, a_1) := \sum_i \lambda_i \langle a_i, a_1 \rangle^2 a_i = \lambda_1 a_1$.
- a_i are **eigenvectors** of the tensor M_3 .

Convert to an orthogonal tensor using pairwise moments

- $M_2 := \mathbb{E}[x_1 \otimes x_2] = \sum_i \lambda_i a_i^{\otimes 2}$.
- Whitening matrix: $W^\top M_2 W = I$.
- Consider tensor $M_3(W, W, W) := \sum_i \lambda_i (W^\top a_i)^{\otimes 3}$. It is an orthogonal tensor.

Parameter Learning in Latent Trees




Learning through Hierarchical Tensor Decomposition

- Assume known tree structure.
- Decompose different triplets: hidden variable is join point on tree.

Alignment issue

- Tensor decomposition is an unsupervised method.
- Hidden labels permuted across different triplets.
- Solution: Align using common node in triplets.

Parameter Learning in Latent Trees




Learning through Hierarchical Tensor Decomposition

- Assume known tree structure.
- Decompose different triplets: hidden variable is join point on tree.

Alignment issue

- Tensor decomposition is an unsupervised method.
- Hidden labels permuted across different triplets.
- Solution: Align using common node in triplets.

Parameter Learning in Latent Trees




Learning through Hierarchical Tensor Decomposition

- Assume known tree structure.
- Decompose different triplets: hidden variable is join point on tree.

Alignment issue

- Tensor decomposition is an unsupervised method.
- Hidden labels permuted across different triplets.
- Solution: Align using common node in triplets.

Parameter Learning in Latent Trees



Learning through Hierarchical Tensor Decomposition

- Assume known tree structure.
- Decompose different triplets: hidden variable is join point on tree.

Alignment issue

- Tensor decomposition is an unsupervised method.
- Hidden labels permuted across different triplets.
- Solution: Align using common node in triplets.

Outline

- 1 Introduction
- 2 Tests for Structure Learning
- 3 Parameter Learning through Tensor Methods
- 4 Integrating Structure and Parameter Learning**
- 5 Conclusion

Integrated Learning

So far..

- Consistent structure learning through sibling tests on distances.
- Parameter learning through tensor decomposition on triplets.

Challenges

- How to integrate structure and parameter learning?
- Can we save on computations through integration?
- Can we learn parameters as we learn the structure?
- Can we **parallelize** learning for scalability?

Integrated Learning

So far..

- Consistent structure learning through sibling tests on distances.
- Parameter learning through tensor decomposition on triplets.

Challenges








- How to integrate structure and parameter learning?
- Can we save on computations through integration?
- Can we learn parameters as we learn the structure?
- Can we **parallelize** learning for scalability?

Key Ideas

- Divide and conquer: find (overlapping) groups of observed variables.
- Learn local subtrees (and parameters) over the groups independently.
- Merge subtrees and tweak parameters to obtain global latent tree model.

Parallel Chow-Liu Based Grouping Algorithm


Minimum spanning tree using information distance $[\hat{d}_{i,j}]$.




Alignment of Parameters

Alignment Correction

- In-group
- Across-group
- Across-neighborhood



Across-neighborhood



Consistency Guarantees

Theorem

The proposed method consistently recovers the structure with $O(\log p)$ samples and parameters with $\text{poly}(p)$ samples.

Extent of parallelism

- Size of groups $\Gamma \leq \Delta^{1+\frac{u}{l}\delta}$.
- Effective depth $\delta := \max_i \{\min_j \{\text{path}(v_i, v_j; \mathcal{T})\}\}$.
- Maximum degree in latent tree: Δ .
- Upper and lower bound on distances between neighbors in the latent tree: u and l .

Implications

- For homogeneous **HMM**, constant sized groups.
- Worst case: star graphs.

Computational Complexity

- N samples, d dimensional observed variables, k state hidden variables.
- p number of observed variables. z non-zero entries per sample.
- Γ sized groups.

Algorithm Steps	Time/worker	Degree of parallelism
Information Distance Estimation	$O(Nz + d + k^3)$	$O(p^2)$
Structure: Minimum Spanning Tree	$O(\log p)$	$O(p^2)$
Structure: Local Recursive Grouping	$O(\Gamma^3)$	$O(p/\Gamma)$
Parameter: Tensor Decomposition	$O(\Gamma k^3 + \Gamma dk^2)$	$O(p/\Gamma)$
Merging and Alignment Correction	$O(dk^2)$	$O(p/\Gamma)$

“Integrated Structure and Parameter Learning in Latent Tree Graphical Models” by F. Huang, U. N. Niranjan, A. Anandkumar. Preprint, June 2014.

Proof Ideas



Relating Chow-Liu Tree with Latent Tree

- Surrogate $\text{Sg}(i)$ for node i : observed node with strongest correlation

$$\text{Sg}(i) := \underset{j \in V}{\operatorname{argmin}} d_{i,j}$$

- Neighborhood preservation

$$(i, j) \in T \Rightarrow (\text{Sg}(i), \text{Sg}(j)) \in T_{\text{ML}}.$$




Chow-Liu grouping reverses edge contractions

Proof by induction

Experiments

- $d = k = 2$ dimensions, $p = 9$ number of variables.



d	p	N	Struct Error	Param Error	Running Time(s)
10	9	50K	0	0.0104	3.8
100	9	50K	0	0.0967	4.4
1000	9	50K	0	0.1014	5.1
10,000	9	50K	0	0.0917	29.9
100,000	9	50k	0	0.0812	56.5
100	9	50K	0	0.0967	10.9
100	81	50K	0.06	0.1814	323.7
100	729	50K	0.16	0.1913	4220.1


Outline

- 1 Introduction
- 2 Tests for Structure Learning
- 3 Parameter Learning through Tensor Methods
- 4 Integrating Structure and Parameter Learning
- 5 Conclusion**

Summary and Outlook

Learning Latent Tree Models

- Integrated Structure and Parameter Learning
- High level of parallelism without losing consistency.



Learning Graphical Model Mixtures

- Tree mixture approximations
- Combinatorial search + spectral decomposition
- Computational and sample guarantees

