Overview of Machine Learning Research

Anima Anandkumar

U.C. Irvine

Learning with Big Data

High Dimensional Regime

- Missing observations, gross corruptions, outliers, ill-posed problems.
- Needle in a haystack: finding low dimensional structures in high dimensional data.

Principled approaches for finding low dimensional structures?

Application 1: Clustering

- Basic operation of grouping data points.
- Hypothesis: each data point belongs to an unknown group.

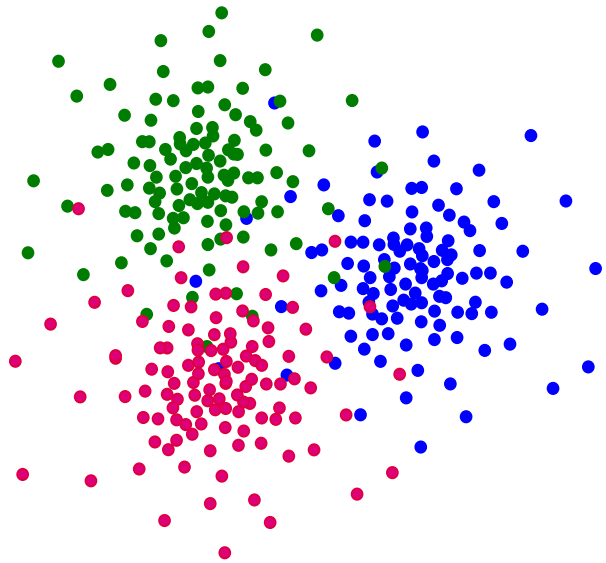

Application 1: Clustering

- Basic operation of grouping data points.
- Hypothesis: each data point belongs to an unknown group.

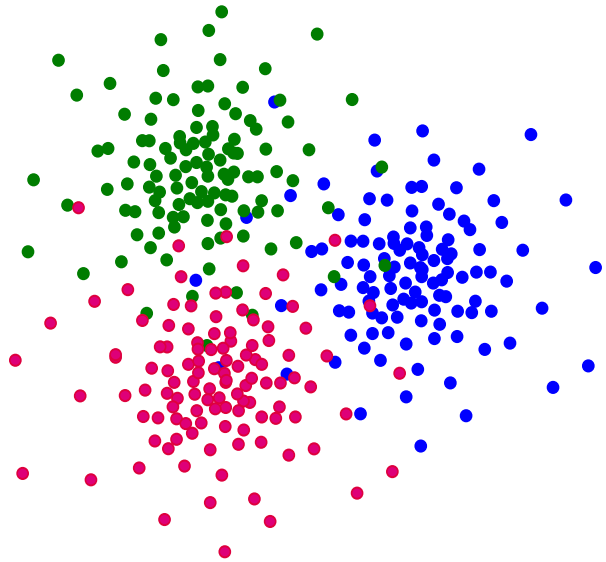

Probabilistic/latent variable viewpoint

- The groups represent different distributions. (e.g. Gaussian).
- Each data point is drawn from one of the given distributions. (e.g. Gaussian mixtures).

Application 2: Topic Modeling

Document modeling

- Observed: words in document corpus.
- Hidden: topics.
- Goal: carry out document summarization.

Application 3: Understanding Human Communities

Social Networks

- Observed: network of social ties, e.g. friendships, co-authorships
- Hidden: groups/communities of actors.

Application 4: Recommender Systems

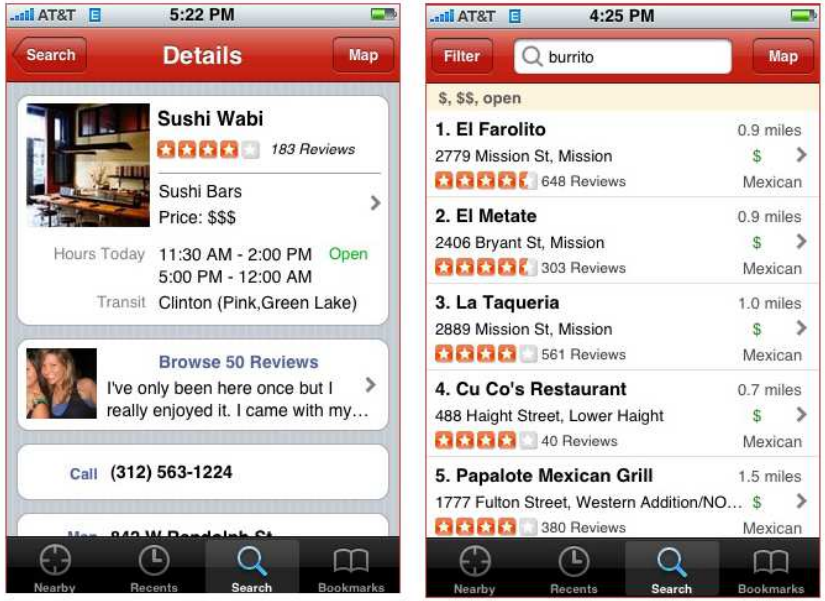

Recommender System

- Observed: Ratings of users for various products, e.g. yelp reviews.
- Goal: Predict new recommendations.
- Modeling: Find groups/communities of users and products.

Application 5: Feature Learning

Label	Features			
0	2.1	5.2	0	0
1	0	0	2	1
1	1.1	0	0	0
0	0	0	7	0
⋮	⋮	⋮	⋮	⋮

Feature Engineering

- Learn good features/representations for classification tasks, e.g. image and speech recognition.
- Sparse representations, low dimensional hidden structures.

Application 6: Computational Biology

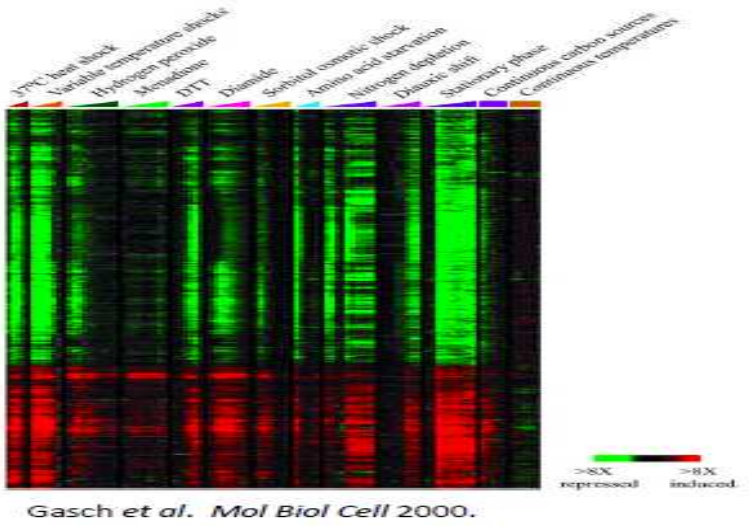

Gasch *et al. Mol Biol Cell* 2000.

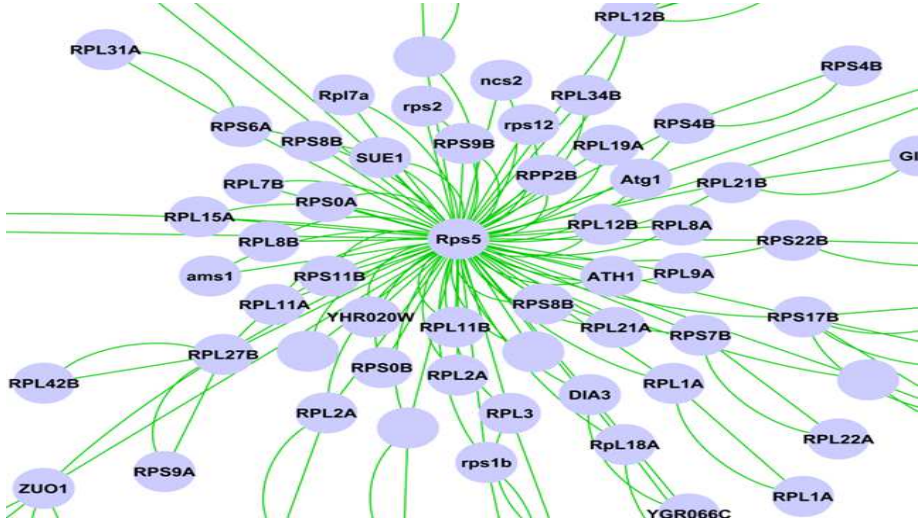

- Observed: gene expression levels
- Goal: discover gene groups
- Hidden variables: regulators controlling gene groups

How to model hidden effects?

Basic Approach: mixtures/clusters

- Hidden variable h is categorical.

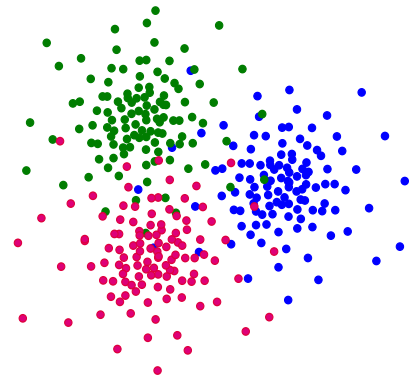

Advanced: Probabilistic models

- Hidden variable h has more general distributions.
- Can model mixed memberships.

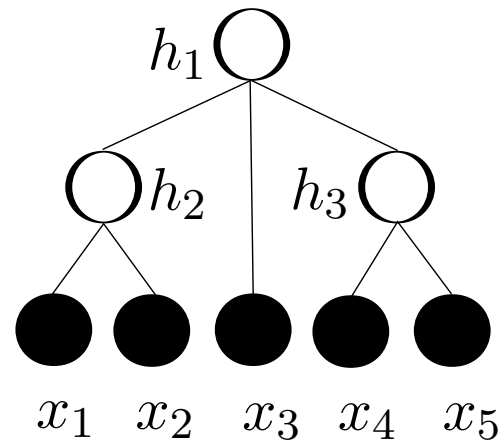

Learning Algorithms through Tensor Factorization

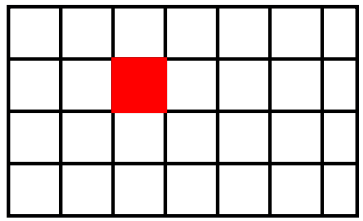

vs.

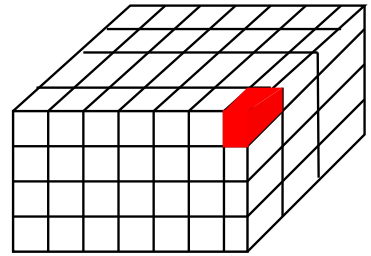

- Co-occurrence of three-words in a document, e.g. [apple, orange, banana].

Tensor Eigenvectors

- Can learn the hidden topics by finding tensor eigenvectors.
- Common friends (neighbors) of triplets of nodes in a social networks.

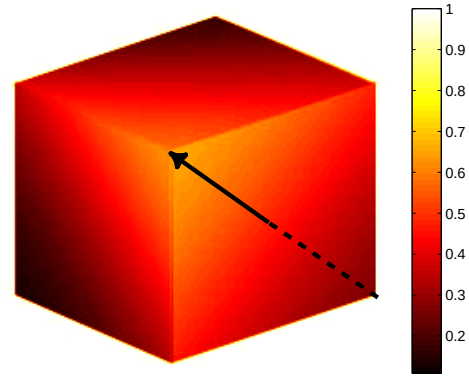

My Research Group

Furong Huang

Majid Janzamin

Hanie Sedghi

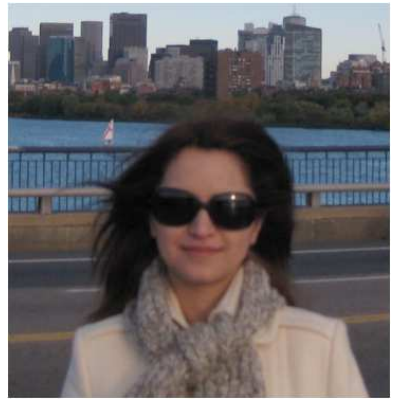

Niranjan UN

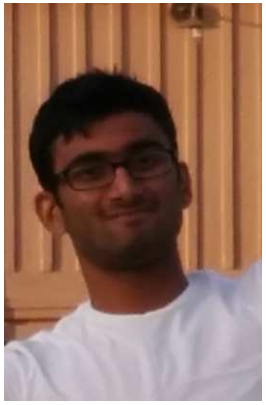

Forough Arabshahi

Resources and Course Information

Resources and Course Information

- ML summer school lectures available at `http://newport.eecs.uci.edu/anandkumar/MLSS.html`
- Publications at `http://newport.eecs.uci.edu/anandkumar/`

Resources and Course Information

- ML summer school lectures available at `http://newport.eecs.uci.edu/anandkumar/MLSS.html`
- Publications at `http://newport.eecs.uci.edu/anandkumar/`

Courses

- EECS 298: Large scale ML: theory and practice.
  - Cloud-based programming, spectral and tensor methods. Hadoop framework.
- EECS 298 (formerly 251B): Statistical learning theory
  - Theoretical course. Non-parametrics, optimization, regularization, concentration bounds.